

## Task-5:

### Sales Prediction Using Python:

Sales prediction means predicting how much of a product people will buy based on factors, a product and service-based business always need their Data Scientist to predict their future sales with every step they take to manipulate the cost of advertising their product.

```
In [1]: 1 #Import required libraries
        2
        3 import numpy as np
        4 import pandas as pd
        5
        6 from sklearn.linear_model import LinearRegression
        7 from sklearn.model_selection import train_test_split
        8 from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_e
        9
        10 import seaborn as sns
        11 import matplotlib.pyplot as plt
```

```
In [2]: 1 df = pd.read_csv("Advertising.csv")
        2 df
```

```
Out[2]:
```

|     | Unnamed: 0 | TV    | Radio | Newspaper | Sales |
|-----|------------|-------|-------|-----------|-------|
| 0   | 1          | 230.1 | 37.8  | 69.2      | 22.1  |
| 1   | 2          | 44.5  | 39.3  | 45.1      | 10.4  |
| 2   | 3          | 17.2  | 45.9  | 69.3      | 9.3   |
| 3   | 4          | 151.5 | 41.3  | 58.5      | 18.5  |
| 4   | 5          | 180.8 | 10.8  | 58.4      | 12.9  |
| ... | ...        | ...   | ...   | ...       | ...   |
| 195 | 196        | 38.2  | 3.7   | 13.8      | 7.6   |
| 196 | 197        | 94.2  | 4.9   | 8.1       | 9.7   |
| 197 | 198        | 177.0 | 9.3   | 6.4       | 12.8  |
| 198 | 199        | 283.6 | 42.0  | 66.2      | 25.5  |
| 199 | 200        | 232.1 | 8.6   | 8.7       | 13.4  |

200 rows × 5 columns

```
In [3]: 1 df.shape
```

```
Out[3]: (200, 5)
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['Unnamed: 0', 'TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')
```

In [5]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      200 non-null   int64
1   TV              200 non-null   float64
2   Radio           200 non-null   float64
3   Newspaper       200 non-null   float64
4   Sales           200 non-null   float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
```

In [6]: 1 *#Checking gthe numbeer of missing values*  
2  
3 df.isnull().sum()

Out[6]: Unnamed: 0 0  
TV 0  
Radio 0  
Newspaper 0  
Sales 0  
dtype: int64

In [7]: 1 df.describe()

Out[7]:

|              | Unnamed: 0 | TV         | Radio      | Newspaper  | Sales      |
|--------------|------------|------------|------------|------------|------------|
| <b>count</b> | 200.000000 | 200.000000 | 200.000000 | 200.000000 | 200.000000 |
| <b>mean</b>  | 100.500000 | 147.042500 | 23.264000  | 30.554000  | 14.022500  |
| <b>std</b>   | 57.879185  | 85.854236  | 14.846809  | 21.778621  | 5.217457   |
| <b>min</b>   | 1.000000   | 0.700000   | 0.000000   | 0.300000   | 1.600000   |
| <b>25%</b>   | 50.750000  | 74.375000  | 9.975000   | 12.750000  | 10.375000  |
| <b>50%</b>   | 100.500000 | 149.750000 | 22.900000  | 25.750000  | 12.900000  |
| <b>75%</b>   | 150.250000 | 218.825000 | 36.525000  | 45.100000  | 17.400000  |
| <b>max</b>   | 200.000000 | 296.400000 | 49.600000  | 114.000000 | 27.000000  |

In [8]: 1 df.corr()

Out[8]:

|                   | Unnamed: 0 | TV       | Radio     | Newspaper | Sales     |
|-------------------|------------|----------|-----------|-----------|-----------|
| <b>Unnamed: 0</b> | 1.000000   | 0.017715 | -0.110680 | -0.154944 | -0.051616 |
| <b>TV</b>         | 0.017715   | 1.000000 | 0.054809  | 0.056648  | 0.782224  |
| <b>Radio</b>      | -0.110680  | 0.054809 | 1.000000  | 0.354104  | 0.576223  |
| <b>Newspaper</b>  | -0.154944  | 0.056648 | 0.354104  | 1.000000  | 0.228299  |
| <b>Sales</b>      | -0.051616  | 0.782224 | 0.576223  | 0.228299  | 1.000000  |

```
In [9]: 1 df.drop("Unnamed: 0",axis=1,inplace=True)
```

```
In [10]: 1 x=df.drop("Sales",axis=1)
        2 y=df["Sales"]
```

```
In [11]: 1 x
```

Out[11]:

|     | TV    | Radio | Newspaper |
|-----|-------|-------|-----------|
| 0   | 230.1 | 37.8  | 69.2      |
| 1   | 44.5  | 39.3  | 45.1      |
| 2   | 17.2  | 45.9  | 69.3      |
| 3   | 151.5 | 41.3  | 58.5      |
| 4   | 180.8 | 10.8  | 58.4      |
| ... | ...   | ...   | ...       |
| 195 | 38.2  | 3.7   | 13.8      |
| 196 | 94.2  | 4.9   | 8.1       |
| 197 | 177.0 | 9.3   | 6.4       |
| 198 | 283.6 | 42.0  | 66.2      |
| 199 | 232.1 | 8.6   | 8.7       |

200 rows × 3 columns

```
In [12]: 1 y
```

Out[12]:

|     |      |
|-----|------|
| 0   | 22.1 |
| 1   | 10.4 |
| 2   | 9.3  |
| 3   | 18.5 |
| 4   | 12.9 |
| ... | ...  |
| 195 | 7.6  |
| 196 | 9.7  |
| 197 | 12.8 |
| 198 | 25.5 |
| 199 | 13.4 |

Name: Sales, Length: 200, dtype: float64

```
In [13]: 1 x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,randc
```

```
In [14]: 1 x_train
```

```
Out[14]:
```

|            | TV    | Radio | Newspaper |
|------------|-------|-------|-----------|
| <b>159</b> | 131.7 | 18.4  | 34.6      |
| <b>62</b>  | 239.3 | 15.5  | 27.3      |
| <b>166</b> | 17.9  | 37.6  | 21.6      |
| <b>45</b>  | 175.1 | 22.5  | 31.5      |
| <b>23</b>  | 228.3 | 16.9  | 26.2      |
| ...        | ...   | ...   | ...       |
| <b>73</b>  | 129.4 | 5.7   | 31.3      |
| <b>144</b> | 96.2  | 14.8  | 38.9      |
| <b>118</b> | 125.7 | 36.9  | 79.2      |
| <b>189</b> | 18.7  | 12.1  | 23.4      |
| <b>99</b>  | 135.2 | 41.7  | 45.9      |

160 rows × 3 columns

```
In [15]: 1 model=LinearRegression()  
2 model
```

```
Out[15]: ▾ LinearRegression  
LinearRegression()
```

```
In [16]: 1 model.fit(x_train,y_train)
```

```
Out[16]: ▾ LinearRegression  
LinearRegression()
```

```
In [17]: 1 y_pred_test=model.predict(x_test)  
2 y_pred_test
```

```
Out[17]: array([ 6.72296139, 13.75631244, 12.04517877,  9.69967488, 15.5690964 ,  
10.05604087, 13.84462636, 11.57199727, 19.50132767, 10.63532913,  
18.86878088,  7.66094904, 14.95238223, 16.44076381, 21.16915592,  
20.34689481, 17.4717922 , 20.44501887,  6.6345129 , 11.59504468,  
10.73339912, 14.9550131 , 20.47462735, 14.4053483 , 19.82336217,  
21.23329569, 13.95616746, 23.25585899,  8.33169334, 20.72918361,  
15.86252542, 15.18601062, 21.74577286, 18.34562985, 12.80653699,  
 9.38349318,  5.37438571, 12.02365628, 19.26275296,  8.12044348])
```

```
In [18]: 1 y_test
```

```
Out[18]: 119      6.6
          77      14.2
          148     10.9
          149     10.1
          154     15.6
          151     11.6
          122     11.6
           6      11.8
          28      18.9
          71      12.4
          188     15.9
          34       9.5
          21      12.5
          40      16.6
          55      23.7
          104     20.7
           3      18.5
          39      21.5
          117      9.4
          134     10.8
          180     10.5
          26      15.0
          54      20.2
          165     11.9
          53      21.2
          93      22.2
          174     11.5
          17      24.4
          56       5.5
          84      21.7
          25      12.0
          156     15.3
          42      20.7
          141     19.2
          50      11.4
          51      10.7
          195      7.6
          116     12.2
          142     20.1
          24       9.7
```

Name: Sales, dtype: float64

```
In [19]: 1 MSE = mean_squared_error(y_test,y_pred_test)
          2 print("Mean squared is : ", MSE)
```

Mean squared is : 2.5509646159531063

```
In [20]: 1 MSE = mean_squared_error(y_pred_test,y_test)
          2 print("Mean squared is : ", MSE)
```

Mean squared is : 2.5509646159531063

```
In [21]: 1 RMSE = np.sqrt(MSE)
2 print("Root mean squared error : ", RMSE)
```

Root mean squared error : 1.5971739466799182

```
In [22]: 1 mean_absolute_error(y_pred_test,y_test)
```

Out[22]: 1.2673937159929234

```
In [23]: 1 r2score=r2_score(y_pred_test,y_test)
2 r2score
```

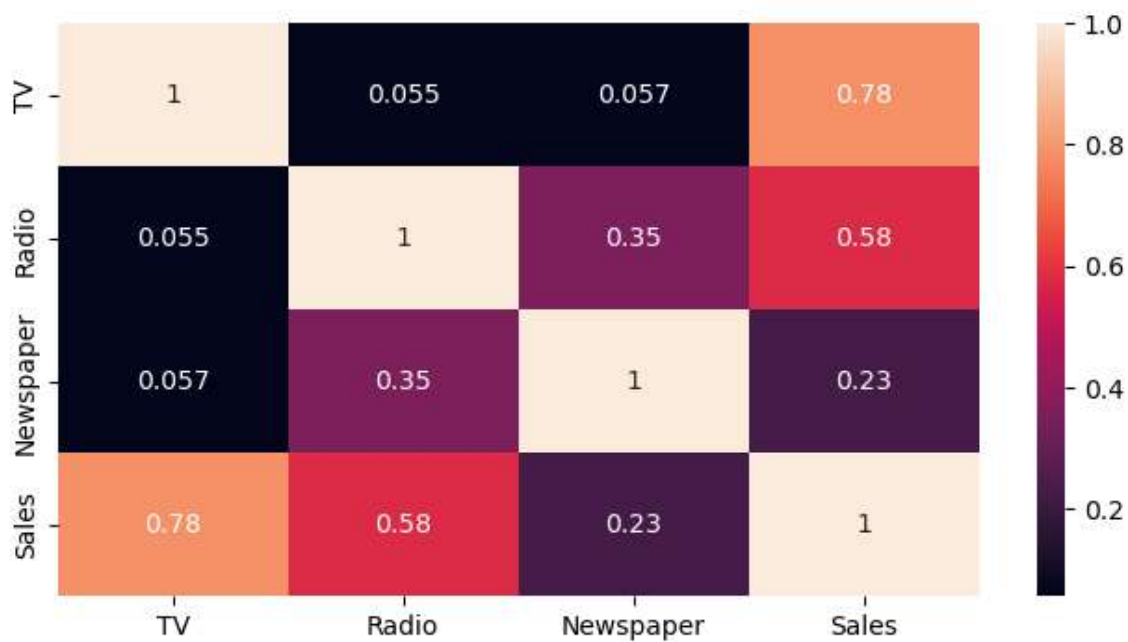
Out[23]: 0.8927421354788028

```
In [24]: 1 r2score=r2_score(y_test,y_pred_test)
2 r2score
```

Out[24]: 0.8984204533332627

```
In [25]: 1 plt.figure(figsize=(8,4))
2 sns.heatmap(df.corr(),annot=True)
```

Out[25]: <Axes: >



```
In [ ]: 1
```